

# LOG AGGREGATION

Log aggregation is the process of **collecting, centralizing, and organizing log data** from many different sources (servers, apps, devices) into **one single location for easier analysis, monitoring, and troubleshooting**.

Telemetry is the automated process of collecting, transmitting, and analyzing data from remote or inaccessible sources to monitor and control systems, providing real-time insights into performance, status, and behavior without physical presence.

**There are two major approaches for achieving this purpose.**

## **Log Forwarding (Push Model) – SaaS model**

This is the most common and robust method. The application or system where the source runs is configured to **push** its logs to a separate, centralized logging server or service.

## **Retrieval (Pull Model) – self hosted**

In this approach, the logs are stored locally (or in a vendor-specific storage), and a separate client or script **pulls** them when needed, typically via a dedicated logging API/agent.

In here we are going to discuss about how to manage and record server logs (container logs) from a centralized location. We proceed with push model with minimal infrastructure configuration. In the other method is able to store logs into local system, setup agent for push logs to the centralized server or else logging agent install itself within the local/remote server (as standalone applications or docker images)

**Our plan is**

- > Structured application logs with Winston (for NodeJS API)
- > Setup Grafana cloud account (cloud is able to act as a SaaS for log management)
- > Setup Grafana Loki as a data source
- > Install Docker plugin that will read logs from Docker containers and ship them to Loki.
- > Configure Docker Deamon. Json for sending logs
- > Inspect logs in Grafana dashboard

Node.js App  
↓ (JSON logs)

Docker stdout  
↓

Promtail (log agent) (in our case log agent is not necessary and configured with docker plugin)

↓

Loki (log database)  
↓

Grafana (UI for team)

## Structured application logs with Winston (for NodeJS API)

One of the best practices is avoiding using plain console logs and moving to structured and leveled comprehensive and detailed application logs. Winston is one of the better libraries for managing structured logs and many more. Additionally, Winston Transporters can be used to ship logs to remote destinations as well.

**npm install winston**

Example logger object with winston (src/logger.index.js)

```
const Winston = require('winston');
const Logtail = require('@logtail/node');
const LogtailTransport = require('winston-transport-logtail');
const LokiTransport = require('winston-loki');
const Winston = require('winston');
const environmentConfig = require('./environmentConfig');
const environmentTypes = require('./environmentConstants/commonConstants');

const customFormat = Winston.format.printf(
  ({level, message, timestamp, stack, ...meta}) => {
    const metaString = Object.keys(meta).length
      ? JSON.stringify(meta, null, 2)
      : '';
    return `${timestamp} | ${level} | ${stack || message} | ${metaString}`;
  }
);

const dev_logger = () => {
  return Winston.createLogger({
    level: 'info',
    format: Winston.format.combine(
      Winston.format.timestamp(),
      Winston.format.timestamp(),
      customFormat,
      // Winston.format.errors({ stack: true }),
      // Winston.format.json()
    ),
    transports: [
      new Winston.transports.Console(), // for Docker stdout
    ],
  });
};

const prod_logger = () => {
  return Winston.createLogger({
    level: 'info',
    format: Winston.format.combine(
      Winston.format.colorize(),
      Winston.format.timestamp(),
      customFormat,
      // Winston.format.errors({ stack: true }),
      // Winston.format.json()
    ),
    transports: [
      new Winston.transports.Console(), // for Docker stdout
    ],
  });
};

// /NOTE: injecting respective logger based on the current environment
let logger = dev_logger();

if(environment.ENVIRONMENT === environmentTypes.PROD){
  logger = prod_logger();
}

module.exports = logger;
```

```
const logger = require('./logger');

logger.info('Server started', { port: 5005 });
logger.error('DB connection failed', { error });
```

```
/*NOTE-----COMMENT-----*/
// const Winston = require('winston');
// const Logtail = require('@logtail/node');
// const LogtailTransport = require('winston-transport-logtail');
// const LokiTransport = require('winston-loki');
// const Winston = require('winston');
// const environmentConfig = require('./environmentConfig');
// const environmentTypes = require('./environmentConstants/commonConstants');

// const customFormat = Winston.format.printf(
//   ({level, message, timestamp, stack, ...meta}) => {
//     const metaString = Object.keys(meta).length
//       ? JSON.stringify(meta, null, 2)
//       : '';
//     return `${timestamp} | ${level} | ${stack || message} | ${metaString}`;
//   }
// );

// const dev_logger = () => {
//   return Winston.createLogger({
//     level: 'info',
//     format: Winston.format.combine(
//       Winston.format.timestamp(),
//       Winston.format.timestamp(),
//       customFormat,
//       // Winston.format.errors({ stack: true }),
//       // Winston.format.json()
//     ),
//     transports: [
//       new Winston.transports.Console(), // for Docker stdout
//     ],
//   });
// };

// const prod_logger = () => {
//   return Winston.createLogger({
//     level: 'info',
//     format: Winston.format.combine(
//       Winston.format.colorize(),
//       Winston.format.timestamp(),
//       customFormat,
//       // Winston.format.errors({ stack: true }),
//       // Winston.format.json()
//     ),
//     transports: [
//       new Winston.transports.Console(), // for Docker stdout
//     ],
//   });
// };

// /NOTE: injecting respective logger based on the current environment
// let logger = dev_logger();

// if(environment.ENVIRONMENT === environmentTypes.PROD){
//   logger = prod_logger();
// }

// module.exports = logger;
```

```
/*NOTE-----COMMENT-----*/
// const Winston = require('winston');
// const LogtailTransport = require('winston-transport-logtail');
// const LokiTransport = require('winston-loki');
// const Winston = require('winston');
// const environmentConfig = require('./environmentConfig');
// const environmentTypes = require('./environmentConstants/commonConstants');

// const customFormat = Winston.format.printf(
//   ({level, message, timestamp, stack, ...meta}) => {
//     const metaString = Object.keys(meta).length
//       ? JSON.stringify(meta, null, 2)
//       : '';
//     return `${timestamp} | ${level} | ${stack || message} | ${metaString}`;
//   }
// );

// const dev_logger = () => {
//   const logger = Winston.createLogger({
//     level: 'info',
//     format: Winston.format.combine(
//       Winston.format.timestamp(),
//       Winston.format.timestamp(),
//       customFormat,
//       // Winston.format.errors({ stack: true }),
//       // Winston.format.json()
//     ),
//     transports: [
//       new Winston.transports.Console(), // for Docker stdout
//     ],
//   });
//   return logger;
// };

// const prod_logger = () => {
//   const logger = Winston.createLogger({
//     level: 'info',
//     format: Winston.format.combine(
//       Winston.format.colorize(),
//       Winston.format.timestamp(),
//       customFormat,
//       // Winston.format.errors({ stack: true }),
//       // Winston.format.json()
//     ),
//     transports: [
//       new Winston.transports.Console(), // for Docker stdout
//     ],
//   });
//   return logger;
// };

// /NOTE: logtail transport is required for local debugging
// if(logger.level === 'error' || logger.level === 'warn'){
//   logger.add(new Winston.transports.Console());
// }

// if(environment.ENVIRONMENT === environmentTypes.PROD){
//   logger.add(new Winston.transports.Logtail({
//     host: environmentConfig.GRAFANA_LOKI_ENDPOINT,
//     token: environmentConfig.GRAFANA_LOKI_TOKEN,
//     environment: environmentConfig.GRAFANA_LOKI_AP_KV,
//     logLevel: 'error',
//     logTimestamp: true,
//     logFormat: 'json',
//     logColor: true
//   }));
// }

// /NOTE: Loki transport (local debugging)
// if(logger.level === 'error' || logger.level === 'warn'){
//   logger.add(new Winston.transports.Loki({
//     host: environmentConfig.GRAFANA_LOKI_ENDPOINT,
//     token: environmentConfig.GRAFANA_LOKI_TOKEN,
//     environment: environmentConfig.GRAFANA_LOKI_AP_KV,
//     logLevel: 'error',
//     logTimestamp: true,
//     logFormat: 'json',
//     logColor: true
//   }));
// }

// /NOTE: Winston Transporter (local debugging)
// if(logger.level === 'error' || logger.level === 'warn'){
//   logger.add(new Winston.transports.Winston({
//     host: environmentConfig.GRAFANA_LOKI_ENDPOINT,
//     token: environmentConfig.GRAFANA_LOKI_TOKEN,
//     environment: environmentConfig.GRAFANA_LOKI_AP_KV,
//     logLevel: 'error',
//     logTimestamp: true,
//     logFormat: 'json',
//     logColor: true
//   }));
// }

// /NOTE: Create logger instance
// if(logger.level === 'error' || logger.level === 'warn'){
//   logger.add(new Winston.transports.Console());
//   logger.add(new Winston.transports.Loki());
//   logger.add(new Winston.transports.Winston());
//   logger.add(new Winston.transports.Logtail());
//   logger.add(new Winston.transports.Loki());
//   logger.add(new Winston.transports.Winston());
// }

// module.exports = logger;
```

## **Setup Grafana cloud account (cloud is able to act as a SaaS for log management)**

<https://grafana.com/cloud>

create free account in Grafana cloud

if you have already an account go with -> My Account

This will display your current cloud stacks

-----  
If you are planning to create self-hosted your own data collector you should create data source manually

connections -> data sources -> add new data source (then configure and test your data source)

But Grafana cloud has already created default loki data source for our use ready for use.  
We can verify in here. connections -> data sources

Now lets see how to obtain loki url and create api tokens

## **Setup Grafana Loki as a data source**

<https://grafana.com/cloud> -> My account -> See the cloud stacks available

Select your cloud stack and -> details -> find the Loki data source available -> details

This will include the necessary details about Loki data source

go to section

Sending Logs to Grafana Cloud using Promtail -> From a standalone host

copy the loki url

clients:

- url: <https://1437567:<Your Grafana.com API Token>@logs-prod-028.grafana.net/loki/api/v1/push>

Now we should generate api token with correct permissions

<https://grafana.com/cloud> -> My account

Navigate into “Access policies” under “security” in left side bar

Create access policy  
Define scopes with permissions  
logs: write  
Create

Now created access policy will show on the list  
Find the option “token -> add token” in created access policy  
Give it a name and create token, copy one time token value and secure

inject the token into this url  
url: <https://1437567:<Your Grafana.com API Token>@logs-prod-028.grafana.net/loki/api/v1/push>

## **Install Docker plugin that will read logs from Docker containers and ship them to Loki.**

<https://grafana.com/docs/loki/latest/send-data/docker-driver>

The Docker plugin must be installed on each Docker host that will be running containers you want to collect logs from.

Run the following command to install the plugin, updating the release version, or changing the architecture (arm64 and amd64 are currently supported), if needed

`docker plugin install grafana/loki-docker-driver:3.3.2-arm64 --alias loki --grant-all-permissions`  
change arm -> amd if needed

`docker plugin ls`

`sudo vim /etc/docker/daemon.json`

place this inside the file. save and exit

## **Configure Docker Deamon. Json for sending logs**

`sudo vim /etc/docker/daemon.json`

place this inside the file. save and exit

```
{  
    "log-driver": "loki",  
    "log-opt": {  
        "loki-url": "https://1437567:<token>logs-prod-028.grafana.net/loki/api/v1/push"  
    }  
}
```

`sudo systemctl restart docker`

## **Inspect logs in Grafana dashboard**

Go to Grafana cloud dashboard

Left side bar -> explore

Left side bar -> drilldown -> logs

query logs

filter logs

select containers and filter specific log

filter based on time range

and many more

**This is only a most fundamental approach for the related topic.**

**Many more things to explore.**